

Parallel Programming in Python: High-level Oomph

Jacob Wilkins

Scientific Computing Department, STFC



Science and
Technology
Facilities Council

August 30, 2023

Overview

- 1 Introduction
- 2 Basic Optimisation
- 3 Vectorisation
- 4 Threading
- 5 ASync
- 6 Multiprocessing
- 7 MPI4Py
- 8 Tensor Computational Graph Libraries
- 9 Summary

Introduction

Questions, questions, questions

- What is Python?

Questions, questions, questions

- What is Python?
- What is parallel?

Questions, questions, questions

- What is Python?
- What is parallel?
- How does it work?

Questions, questions, questions

- What is Python?
- What is parallel?
- How does it work?
- How to make the most of it.

Questions, questions, questions

- What is Python?
- What is parallel?
- How does it work?
- How to make the most of it.
- Not intended as introduction to parallel programming!

- Python is a high-level interpreted programming language

- Python is a high-level interpreted programming language
- Most versions run through C (CPython, not Cython)

- Python is a high-level interpreted programming language
- Most versions run through C (CPython, not Cython)
- In general, slow!

- Python is a high-level interpreted programming language
- Most versions run through C (CPython, not Cython)
- In general, slow!
- Use libraries to get real performance.

Snakes are fish now

- Python has a GIL

Snakes are fish now

- Python has a GIL
- Global-interpreter lock

Snakes are fish now

- Python has a GIL
- Global-interpreter lock
- No Python process can desynchronise from any other


Snakes are fish now

- Python has a GIL
- Global-interpreter lock
- No Python process can desynchronise from any other
- Efforts to remove GIL have been ongoing for years

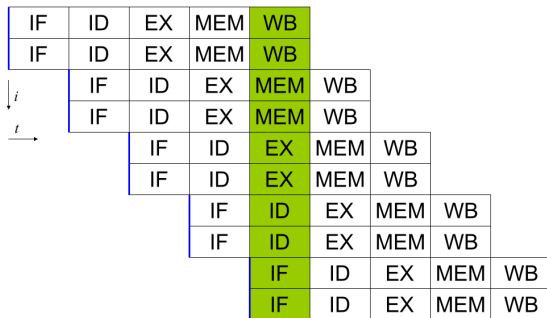
Snakes are fish now

- Python has a GIL
- Global-interpretor lock
- No Python process can desynchronise from any other
- Efforts to remove GIL have been ongoing for years
- May finally be reaching fruition

- What does it mean to do jobs in parallel?




- What does it mean to do jobs in parallel?




Fork-Join

- Fork-Join parallelism is temporary parallelism.




Fork-Join

- Fork-Join parallelism is temporary parallelism.
- Takes time to start/stop threads (sometimes leave them running).




Fork-Join

- Fork-Join parallelism is temporary parallelism.
- Takes time to start/stop threads (sometimes leave them running).
- Often simpler to deal with for small regions.



Fork-Join

- Fork-Join parallelism is temporary parallelism.
- Takes time to start/stop threads (sometimes leave them running).
- Often simpler to deal with for small regions.
- Often less efficient.



- Independent parallelism has full independent threads.

Independent

- Independent parallelism has full independent threads.
- Always parallel, need ground-up supporting structure.

- Independent parallelism has full independent threads.
- Always parallel, need ground-up supporting structure.
- Not in sync, running own copies of codes.

- Independent parallelism has full independent threads.
- Always parallel, need ground-up supporting structure.
- Not in sync, running own copies of codes.
- Can be periodically forcibly synchronised.

- All processes see the same memory

Shared Memory

- All processes see the same memory
- All processes can change the same memory

Shared Memory

- All processes see the same memory
- All processes can change the same memory
- Race conditions are a risk

- All processes see the same memory
- All processes can change the same memory
- Race conditions are a risk
- Race condition is where 2 threads try to access same memory at once

Distributed Memory

- All processors have their own memory

Distributed Memory

- All processors have their own memory
- Need to specify which memory to transfer

- All processors have their own memory
- Need to specify which memory to transfer
- Some systems allow memory windows

Basic Optimisation

Using Generators

- Avoid computing things you aren't going to use

Using Generators

- Avoid computing things you aren't going to use
- Generators allow infinite loops

Using Generators

- Avoid computing things you aren't going to use
- Generators allow infinite loops
- Functions only calculated as needed

Using Generators

- Avoid computing things you aren't going to use
- Generators allow infinite loops
- Functions only calculated as needed
- Maps and filters can apply sequentially

```
def my_simple_counter(n: int = 0):  
    """Count infinitely"""  
    while true:  
        yield n  
        n = n+1  
  
for i in my_simple_counter(15):  
    print(i) # 15, 16, 17, ...
```


- Simple comprehensions run in C

- Simple comprehensions run in C
- Efficient clean use of these can speed up code

- Simple comprehensions run in C
- Efficient clean use of these can speed up code
- Also implement filters and transformations

No comprende

```
# List comprehensions
[3*x for x in range(10) if x % 2]
# Set comprehensions
{3*x for x in range(10) if x % 2}
# Dict comprehensions
{key: 3*val for key, val in zip("abcdefghij",
                                range(10)) if val % 2}
# Generator comprehension
(3*x for x in range(10) if x % 2)
# Using generator comprehensions
tuple(3*x for x in range(10) if x % 2)
"\t".join(str(3*x) for x in range(10) if x % 2)
```

Vectorisation

- Numpy provides free parallelism out of the gate

```
import numpy as np

x = np.fromiter(range(1000))
print(x + np.ones(1000))

y = np.sin(x)
```

- Numpy provides free parallelism out of the gate
- Numpy operations on arrays vectorised in C

```
import numpy as np

x = np.fromiter(range(1000))
print(x + np.ones(1000))

y = np.sin(x)
```

- Numpy provides free parallelism out of the gate
- Numpy operations on arrays vectorised in C
- Numpy ufuncs operate efficiently over arrays

```
import numpy as np
```

```
x = np.fromiter(range(1000))  
print(x + np.ones(1000))
```

```
y = np.sin(x)
```


No, you func

```
import numpy as np
```

```
def add_two(x):  
    "Add 2 to x"  
    return x + 2
```

```
#                               frompyfunc(func, nin, nout)  
my_ufunc = np.frompyfunc(add_two, 1, 1)  
print(my_ufunc(np.array([3, 4, 5]))) # => [5 6 7]
```

Ufunc powers

- `__call__(*args, **kwargs)`
Call self as a function.
- `accumulate(array[, axis, dtype, out])`
Accumulate the result of applying the operator to all elements.
- `at(a, indices[, b])`
Performs unbuffered in place operation on operand 'a' for elements specified by 'indices'.
- `outer(A, B, /, **kwargs)`
Apply the ufunc op to all pairs (a, b) with a in A and b in B.
- `reduce(array[, axis, dtype, out, keepdims, ...])`
Reduces array's dimension by one, by applying ufunc along one axis.
- `reduceat(array, indices[, axis, dtype, out])`
Performs a (local) reduce with specified slices over a single axis.

Don't vectorise

- Don't use `numpy.vectorize`

The `vectorize` function is provided primarily for convenience, not for performance. The implementation is essentially a for loop.

– Numpy Docs

<https://numpy.org/doc/stable/reference/generated/numpy.vectorize.html>

Threading

- Fork-join, shared memory model

Threads in briefs

- Fork-join, shared memory model
- Throttled by the GIL

Threads in briefs

- Fork-join, shared memory model
- Throttled by the GIL
- Mostly designed for background tasks where the main loop shouldn't freeze

Threads in briefs

- Fork-join, shared memory model
- Throttled by the GIL
- Mostly designed for background tasks where the main loop shouldn't freeze
- Think GUIs, IO operations, etc.

Threads in briefs

```
from threading import Thread

# Declare the job and the arguments
p = Thread(target=f, args=('bob',))
# Starts a thread running the job – Fork
p.start()
# Wait until finished
p.join()
```

- In CPython, due to the Global Interpreter Lock, only one thread can execute Python code at once (even though certain performance-oriented libraries might overcome this limitation).
 - Python Docs

<https://docs.python.org/3/library/threading.html>

- More like writing a note to yourself
- Do the job when the processor has free time

- Need to be aware of “thread-safe” methods.
- Queues are useful for managing threads
- Threads are also handy for spawning other processes.

ASync

- In language native co-routines

- In language native co-routines
- Ask now, get later

- In language native co-routines
- Ask now, get later
- Similar to `threading`, but subtly different

- Two main keywords `async`, `await`
- Also involves the `asyncio` package.

```
import asyncio
```

```
async def sleepy():  
    print("Yawn")  
    asyncio.sleep(1)  
    print("Awake")
```

```
# Runs in 1 second
```

```
asyncio.gather(sleepy(), sleepy(), sleepy())
```

```
# Takes all 3
```

```
for _ in range(3):  
    sleepy()
```

- `async` attaches to `def`, `for` and `with`

```
async def g(x):  
    return 2*x
```

```
async def f(x):  
    ans = await g(x)  
    return ans
```

- `async` attaches to `def`, `for` and `with`
- `async` declares that object is to be scheduled in background.

```
async def g(x):  
    return 2*x
```

```
async def f(x):  
    ans = await g(x)  
    return ans
```

- `async` attaches to `def`, `for` and `with`
- `async` declares that object is to be scheduled in background.
- `await` attaches to `async` functions

```
async def g(x):  
    return 2*x
```

```
async def f(x):  
    ans = await g(x)  
    return ans
```

- `async` attaches to `def`, `for` and `with`
- `async` declares that object is to be scheduled in background.
- `await` attaches to `async` functions
- `await` says to pause execution until answer returned

```
async def g(x):  
    return 2*x
```

```
async def f(x):  
    ans = await g(x)  
    return ans
```

- Can help to think in terms of generators
- `async` declares object whose value is not computed yet
- `await` acts like `yield` from for `async`

Async objects

- Can declare object as awaitable
- Add `__await__` method
- Must return iterator object

```
class MyObject:
    def __init__(self):
        ...

    def __await__(self):
        ...
        return iterator_object
```

Async generator

- Like normal generators
- Can be iterated over with `async for`
- **N.B.** this is not parallelism, merely allowing other operations to borrow the processor.

```
import asyncio
```

```
async def async_generator(n=0):  
    while true:  
        yield n  
        n = n + 1  
        asyncio.sleep(1)
```

```
async for i in async_generator(3):  
    print(i)
```


Multiprocessing

Finally Parallel!

- Fork-join, shared memory model

```
from multiprocessing import Process
```

```
def f(name):  
    print('hello ', name)
```

```
# Declare the job and the arguments  
p = Process(target=f, args=('bob',))  
# Starts a thread running the job  
p.start()  
# Wait until finished  
p.join()
```

Finally Parallel!

- Fork-join, shared memory model
- Actually sidesteps the GIL – Finally parallel

```
from multiprocessing import Process
```

```
def f(name):  
    print('hello ', name)
```

```
# Declare the job and the arguments  
p = Process(target=f, args=('bob',))  
# Starts a thread running the job  
p.start()  
# Wait until finished  
p.join()
```

Finally Parallel!

- Fork-join, shared memory model
- Actually sidesteps the GIL – Finally parallel
- Support for task based parallelism

```
from multiprocessing import Process
```

```
def f(name):  
    print('hello ', name)
```

```
# Declare the job and the arguments  
p = Process(target=f, args=('bob',))  
# Starts a thread running the job  
p.start()  
# Wait until finished  
p.join()
```

Multiprocessing shared memory

- Support for shared memory processing

```
from multiprocessing import Process, Value, Array
```

```
def f(n, a):  
    n.value = 3.1415927  
    for i, val in enumerate(a):  
        a[i] = -val
```

```
# Double
```

```
num = Value('d', 0.0)
```

```
# Integer
```

```
arr = Array('i', [4, 5, 6])
```

```
p = Process(target=f, args=(num, arr))
```

```
p.start()
```

```
p.join()
```

Swimming in threads

- Support for iterative parallelism

```
from multiprocessing import Pool
```

```
def add_two(x):  
    return x + 2
```

```
# start 4 worker processes  
with Pool(processes=4) as pool:
```

```
    x = pool.map(add_two, [4, 5, 6]) # 6, 7, 8  
    y = pool.starmap(add, [(1, 2), (3, 4)]) # 3, 7
```

Swimming in threads

- Support for iterative parallelism
- Threads are grouped into pools

```
from multiprocessing import Pool
```

```
def add_two(x):  
    return x + 2
```

```
# start 4 worker processes  
with Pool(processes=4) as pool:
```

```
    x = pool.map(add_two, [4, 5, 6]) # 6, 7, 8  
    y = pool.starmap(add, [(1, 2), (3, 4)]) # 3, 7
```

Swimming in threads

- Support for iterative parallelism
- Threads are grouped into pools
- Simplest way is to parallelise over array

```
from multiprocessing import Pool
```

```
def add_two(x):  
    return x + 2
```

```
# start 4 worker processes  
with Pool(processes=4) as pool:
```

```
    x = pool.map(add_two, [4, 5, 6]) # 6, 7, 8  
    y = pool.starmap(add, [(1, 2), (3, 4)]) # 3, 7
```


MPI4Py

- MPI is a standard for parallel computing

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD
```

```
size = comm.Get_size()
```

```
rank = comm.Get_rank()
```

```
print(f"Hello from process {rank} of {size}")
```

- MPI is a standard for parallel computing
- Originally designed to enable parallel computing in languages such as C and Fortran

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD
```

```
size = comm.Get_size()
```

```
rank = comm.Get_rank()
```

```
print(f"Hello from process {rank} of {size}")
```

- MPI is a standard for parallel computing
- Originally designed to enable parallel computing in languages such as C and Fortran
- Since it exists in C, is available in Python

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD  
size = comm.Get_size()  
rank = comm.Get_rank()
```

```
print(f"Hello from process {rank} of {size}")
```

- MPI is a standard for parallel computing
- Originally designed to enable parallel computing in languages such as C and Fortran
- Since it exists in C, is available in Python
- `mpi4py` is that interface to underlying C

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD  
size = comm.Get_size()  
rank = comm.Get_rank()
```

```
print(f"Hello from process {rank} of {size}")
```

- Independent, distributed memory model

```
>> mpirun -np 4 python parallel_python.py  
Hello from process 3 of 4  
Hello from process 0 of 4  
Hello from process 1 of 4  
Hello from process 2 of 4
```

- Independent, distributed memory model
- Always parallel

```
>> mpirun -np 4 python parallel_python.py  
Hello from process 3 of 4  
Hello from process 0 of 4  
Hello from process 1 of 4  
Hello from process 2 of 4
```

- Independent, distributed memory model
- Always parallel
 - Requires you to change how you think from the ground up

```
>> mpirun -np 4 python parallel_python.py  
Hello from process 3 of 4  
Hello from process 0 of 4  
Hello from process 1 of 4  
Hello from process 2 of 4
```


- Independent, distributed memory model
- Always parallel
 - Requires you to change how you think from the ground up
- Effective up to thousands of cores in principle

```
>> mpirun -np 4 python parallel_python.py  
Hello from process 3 of 4  
Hello from process 0 of 4  
Hello from process 1 of 4  
Hello from process 2 of 4
```

- Independent, distributed memory model
- Always parallel
 - Requires you to change how you think from the ground up
- Effective up to thousands of cores in principle
- Each process has its own memory pool

```
>> mpirun -np 4 python parallel_python.py  
Hello from process 3 of 4  
Hello from process 0 of 4  
Hello from process 1 of 4  
Hello from process 2 of 4
```

- Independent, distributed memory model
- Always parallel
 - Requires you to change how you think from the ground up
- Effective up to thousands of cores in principle
- Each process has its own memory pool
- Need to send all data to other processes manually and keep track yourself.

```
>> mpirun -np 4 python parallel_python.py
Hello from process 3 of 4
Hello from process 0 of 4
Hello from process 1 of 4
Hello from process 2 of 4
```

Communicators

- Fundamental object in MPI4py is the communicator.
- Communicator groups and labels processors.
- Enables transfer of data between them.
- `mpi4py.MPI.COMM_WORLD` is default.

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD
```

```
# Create new communicator for procs (1,2),(3,4)  
colour = int(comm.rank < 3) # Group to be in  
key = comm.rank # Rank order on comm  
new_comm = comm.split(colour, key)
```

```
# Create a 2x2 non-periodic grid of processors  
comm.Create_cart([2,2], periods=[False, False])
```

What's the point(-to-point)?

- Point-to-point communications specify source and destination
- Collective communications specify source
- Collectives send/receive from all in group
- MPI4py compatible with numpy
- Cartesian communicators allow N-D destinations

```
# Send from rank  $\rightarrow$  rank + 1 (next in series)  
comm.send(3, dest=(rank+1)%size)
```

```
# Split array and send equal portions to each  
  processor in group  
recv = np.zeros(50)  
send = np.array(100)  
send[0:50] = 1; send[50:] = 2;  
comm.Scatter(send, recv, root=0)
```

MPI Differences

- Don't need to handle typing explicitly (uses pickle)
- Requests from asynchronous comms are objects
- lowercase methods transfer objects
- Titlecase methods transfer buffers (lower-level)

```
comm.bcast(data, root=0)  
comm.Bcast((data, MPI.DOUBLE), root=0)
```

```
req = comm.isend(data, rank+1)  
req.test()  
req.wait()
```

- Works with latest MPI features:

- Works with latest MPI features:
 - MPI-IO

- Works with latest MPI features:
 - MPI-IO
 - One-sided communications (memory windows)

- Works with latest MPI features:
 - MPI-IO
 - One-sided communications (memory windows)
 - GPU Aware MPI

- Works with latest MPI features:
 - MPI-IO
 - One-sided communications (memory windows)
 - GPU Aware MPI
- Some Python specific:

- Works with latest MPI features:
 - MPI-IO
 - One-sided communications (memory windows)
 - GPU Aware MPI
- Some Python specific:
 - MPI Sub-Process spawning

- Works with latest MPI features:
 - MPI-IO
 - One-sided communications (memory windows)
 - GPU Aware MPI
- Some Python specific:
 - MPI Sub-Process spawning
 - SWIG

- Works with latest MPI features:
 - MPI-IO
 - One-sided communications (memory windows)
 - GPU Aware MPI
- Some Python specific:
 - MPI Sub-Process spawning
 - SWIG
 - F2Py

Tensor Computational Graph Libraries

- Provide high-level interfaces for complicated mathematical operations

- Provide high-level interfaces for complicated mathematical operations
- Will often compile (JIT) into other languages

- Provide high-level interfaces for complicated mathematical operations
- Will often compile (JIT) into other languages
- Often include support for GPUs

- Aesara (previously Theano)

- Aesara (previously Theano)
- PyTorch

- Aesara (previously Theano)
- PyTorch
- JAX(ish)

- Aesara (previously Theano)
- PyTorch
- JAX(ish)
- Tensorflow

- Aesara (previously Theano)
- PyTorch
- JAX(ish)
- Tensorflow
- Going to take a look at Aesara

- Aesara works by defining a sequence of functions on anonymous typed arguments
- Can then dispatch those functions as needed to different backends
- Designed to be generic and able to operate generally


```
import aesara
from aesara import tensor as at


# Declare two symbolic floating-point scalars
a = at.dscalar("a")
b = at.dscalar("b")

# Create a simple example expression
c = a + b

# Convert the expression into a callable object
that takes '(a, b)'
# values as input and computes the value of 'c'.
f_c = aesara.function([a, b], c)

assert f_c(1.5, 2.5) == 4.0
```

- Builds a graph of operations
- Attempts to optimise and reduce graph
- Compiles resultant graph returning accessible function
- Compiled function dispatched to “processor”
- “Processor” may be CPU, GPU, TPU, or more



<https://github.com/aesara-devs/aesara/>

- In principle saves you the work of: Optimisation, Ordering, Dispatching

```

d = a/a + (M + a).dot(v)
aesara.dprint(d)
# Elemwise{add,no_inplace} [id A] ''
# | InplaceDimShuffle{x} [id B] ''
# | | Elemwise{true_divide,no_inplace} [id C] ''
# | | a [id D]
# | | a [id D]
# | dot [id E] ''
# | Elemwise{add,no_inplace} [id F] ''
# | | M [id G]
# | | InplaceDimShuffle{x,x} [id H] ''
# | | a [id D]
# | v [id I]

```

- In principle saves you the work of: Optimisation, Ordering, Dispatching

```
f_d = aesara.function([a, v, M], d)
# 'a/a' → '1' and the dot product is replaced
#   with a BLAS function
# Elemwise{Add}[(0, 1)] [id A] '' 5
# | TensorConstant{(1,) of 1.0} [id B]
# | CGemv{inplace} [id C] '' 4
# | AllocEmpty{dtype='float64'} [id D] '' 3
# | | Shape_i{0} [id E] '' 2
# | | M [id F]
# | TensorConstant{1.0} [id G]
# | Elemwise{add,no_inplace} [id H] '' 1
# | | M [id F]
# | | InplaceDimShuffle{x,x} [id I] '' 0
# | | a [id J]
# | v [id K]
# | TensorConstant{0.0} [id L]
```

- Because these have access to (or are) the graph
- Can manipulate the graph constructively
- Can compute gradients or inverses

```
x = at.dmatrix('x')  
# Logistic equation  
s = at.sum(1 / (1 + at.exp(-x)))  
gs = at.grad(s, x)
```

```
# Now have a function for derivative of logistic  
dlogistic = aesara.function([x], gs)
```

Summary

- Numpy – Parallel(ish)

Summary

- Numpy – Parallel(ish)
- Threading – Not parallel

Summary

- Numpy – Parallel(ish)
- Threading – Not parallel
- Async – Not parallel

Summary

- Numpy – Parallel(ish)
- Threading – Not parallel
- Async – Not parallel
- Multiprocessing – Parallel

Summary

- Numpy – Parallel(ish)
- Threading – Not parallel
- Async – Not parallel
- Multiprocessing – Parallel
- MPI4Py – Parallel

Summary

- Numpy – Parallel(ish)
- Threading – Not parallel
- Async – Not parallel
- Multiprocessing – Parallel
- MPI4Py – Parallel
- TCGM – Maybe parallel

- If you want to have a play there is a really bad Python MD code available at:
https://github.com/oerc0122/worst_md/blob/main/md.py